

MovieLens Project

Author: Gabrielmario Scozzarro ; Date:21/05/2020

Executive summary

GroupLens is a research lab in the Department of Computer Science and Engineering at the University of Minnesota. The GroupLens team has collected and made available rating data sets from the MovieLens web site (<http://movielens.org> (<http://movielens.org>)). The data sets were collected over various periods of time and for this reason it is very suitable for the implementation of machine learning algorithms. Those algorithms improve automatically through data or what we can call "experience". It is considered as a subset of AI. Machine learning algorithms are based on well known mathematical models fed with collected data, known as "training data", in order to make predictions or decisions without being explicitly programmed to do so. Machine learning algorithms are used in a wide variety of applications, spanning from IT to robotics applications, where it is difficult or infeasible to develop conventional algorithms to perform the needed tasks.

The goal of this project is to build a rating recommender system using the machine learning fed with the MovieLens data set. This report will guide you through the methods and the results of this project. For an easy and immediate comparison with other approaches and for quality estimate we will use the residual mean squared error (RMSE) as the main KPI with a target value of 0.8641.

Methods and analysis

We use the MovieLens 10M dataset that consists of 10M observed ratings from almost 70K different users applied to 10K different movies. We already divided the data set in 2 parts with a proportion of 90/10 for train_set and test_set respectively.

In [51]:

```
train_set<- readRDS(file = "edx.rds")
test_set<- readRDS(file = "validation.rds")
```

To perform a preliminary analysis of the dataset we use the following packages:

In [52]:

```
library(tidyverse)
library(lubridate)
```

and options:

In [53]:

```
options(repr.plot.width=18, repr.plot.height=9)
```

The data structure is:

In [54]:

```
dim(train_set)
```

9000055 · 6

In [55]:

```
head(train_set, 10)
```

A data.frame: 10 × 6

	userId	movieId	rating	timestamp	title	genres
	<int>	<dbl>	<dbl>	<int>	<chr>	<chr>
1	1	122	5	838985046	Boomerang (1992)	Comedy Romance
2	1	185	5	838983525	Net, The (1995)	Action Crime Thriller
4	1	292	5	838983421	Outbreak (1995)	Action Drama Sci-Fi Thriller
5	1	316	5	838983392	Stargate (1994)	Action Adventure Sci-Fi
6	1	329	5	838983392	Star Trek: Generations (1994)	Action Adventure Drama Sci-Fi
7	1	355	5	838984474	Flintstones, The (1994)	Children Comedy Fantasy
8	1	356	5	838983653	Forrest Gump (1994)	Comedy Drama Romance War
9	1	362	5	838984885	Jungle Book, The (1994)	Adventure Children Romance
10	1	364	5	838983707	Lion King, The (1994)	Adventure Animation Children Drama Musical
11	1	370	5	838984596	Naked Gun 33 1/3: The Final Insult (1994)	Action Comedy

A deeper analysis on the rating values:

In [56]:

```
summary(train_set$rating)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.500	3.000	4.000	3.512	4.000	5.000

To perform an analysis of the rating behavior through the years, on the train_set we convert the time stamp in DateTime format:


In [57]:

```
train_set<- train_set %>%  
  mutate(timestamp = as_datetime(timestamp))  
  
train_set %>%  
  mutate(year = year(timestamp)) %>%  
  group_by(year) %>%  
  summarize(n_rating = n()) %>%  
  print(n = 15)
```

```
# A tibble: 15 x 2  
  year n_rating  
  <dbl>   <int>  
1  1995         2  
2  1996    942772  
3  1997    414101  
4  1998    181634  
5  1999    709893  
6  2000   1144349  
7  2001    683355  
8  2002    524959  
9  2003    619938  
10 2004    691429  
11 2005   1059277  
12 2006    689315  
13 2007    629168  
14 2008    696740  
15 2009    13123
```

In [58]:

```
train_set %>%  
  mutate(year = year(timestamp)) %>%  
  group_by(year) %>%  
  summarize(n_rating = n()) %>%  
  ggplot(aes(year, n_rating)) +  
  geom_area(fill="#69b3a2", alpha=0.4) +  
  geom_line(color="#69b3a2", size=2) +  
  geom_point(size=3, color="#69b3a2") +  
  ggtitle("Number of rates per year")
```




We can observe that in the train_set we have 3 spikes 1996, 2000 and 2005.

Here's a distribution of the user ratings in the train_set:

In [59]:

```
train_set %>%  
  group_by(rating) %>%  
  summarize(dens = n()) %>%  
  ggplot(aes(rating,dens)) +  
  geom_area( fill="#69b3a2", alpha=0.4) +  
  geom_line(color="#69b3a2", size=2) +  
  geom_point(size=3, color="#69b3a2") +  
  ggtitle("Rating distribution")
```



It appears that users are quite generous in their ratings. In fact, as we seen in a previous analysis the mean rate is 3.51, most of the rates are on the interval 3-5 with a prevalence for round value. We can also investigate how much users are willing to rate a movie like this:

In [60]:

```
length(unique(train_set$userId))
```

69878

In [61]:

```
train_set%>%  
  group_by(userId) %>%  
  summarize(n_rating = n()) %>%  
  summary()
```

	userId	n_rating
Min.	: 1	Min. : 10.0
1st Qu.:	17943	1st Qu.: 32.0
Median :	35799	Median : 62.0
Mean :	35782	Mean : 128.8
3rd Qu.:	53620	3rd Qu.: 141.0
Max.	: 71567	Max. : 6616.0

There are 69878 unic users. Each user rated at least 10 movies, so the distribution should not be caused just by chance variance in the quality of movies.

We can now explore the genres feature like this:

In [62]:

```
train_set%>%  
  group_by(genres) %>%  
  summarize(movies_per_genre = n()) %>%  
  top_n(10, movies_per_genre)
```


A tibble: 10 × 2

genres	movies_per_genre
<chr>	<int>
Action Adventure Sci-Fi	219938
Action Adventure Thriller	149091
Comedy	700889
Comedy Drama	323637
Comedy Drama Romance	261425
Comedy Romance	365468
Crime Drama	137387
Drama	733296
Drama Romance	259355
Drama Thriller	145373

We can see that movies can be categorized using one or more genre and which are the most present genres in the train_set. We can also compute the average and standard error for each category and find the category with the highest or lowest average rating:

In [63]:

```
train_set %>% group_by(genres) %>%
  summarize(n_rating = n(), avg_rating = mean(rating), se_rating = sd(rating)/sqrt(n
  ())) %>%
  filter(n_rating >= 50000) %>%
  mutate(genres = reorder(genres, avg_rating)) %>%
  ggplot(aes(x = genres, y = avg_rating, ymin = avg_rating - 2*se_rating, ymax = avg_ra
  ting + 2*se_rating)) +
  geom_point(color = "#69b3a2" ) +
  geom_errorbar(color = "#69b3a2") +
  theme(axis.text.x = element_text(angle = 90, hjust = 1)) +
  ggtitle("Average and standard error for each category")
```



With a filter that keeps only genres with more than 50K ratings we can appreciate that Horror movies has the lowest average ratings and Comedy|Crime|Drama the highest. Now we can investigate which movie title has the best ratings like this:

In [64]:

```
length(unique(train_set$movieId))
```

10677

In [65]:

```
train_set %>% group_by(movieId, title) %>%
  summarize(count = n()) %>%
  arrange(desc(count)) %>%
  head(10)
```

A grouped_df: 10 × 3

movieId	title	count
<dbl>	<chr>	<int>
296	Pulp Fiction (1994)	31362
356	Forrest Gump (1994)	31079
593	Silence of the Lambs, The (1991)	30382
480	Jurassic Park (1993)	29360
318	Shawshank Redemption, The (1994)	28015
110	Braveheart (1995)	26212
457	Fugitive, The (1993)	25998
589	Terminator 2: Judgment Day (1991)	25984
260	Star Wars: Episode IV - A New Hope (a.k.a. Star Wars) (1977)	25672
150	Apollo 13 (1995)	24284

Among the 10677 different movies rated from 1995 to 2019 these are the top 10 most rated.

Recommendation system

For a recommendation system we need to identify the most important features that can contain helpful information to predict the rating any given user will assign to any movie. We start defining the RMSE formula needed for the quality assessment of the process we are going to build. We choose to apply a collaborative filtering approach with a Matrix Factorization to increment the accuracy of the prediction.

As known collaborative filtering is a method for automatic predictions of users' interests by collecting preferences or taste information from many other users. The reasoning behind this approach is that if a user A has the same sentiment as a user B on a matter, A is more likely to have B's opinion on another matter compared to a randomly chosen person. Note that the most fascinating thing about this approach is that these predictions are specific to the user, but use information is harvested from many users.

To accomplish our goal we will make use of the `recoSystem` package which is an R wrapper of the LIBMF, an open source library for recommender system using parallel matrix factorization. In addition this package allows us to speed up all processes using parallel computation on many threads.

In [78]:

```
library(recoSystem)
```

Defining RMSE:

$$RMSE = \sqrt{\frac{1}{N} \sum_{u,i} (\hat{y}_{u,i} - y_{u,i})^2}$$

In [67]:

```
RMSE<- function(true_ratings, predicted_ratings){
  sqrt(mean((true_ratings - predicted_ratings)^2))
}
```

Now we can convert the 2 data set in the recosystem desired input format:

In [68]:

```
train_data <- with(train_set, data_memory(user_index = userId, item_index = movieId, rating = rating))

test_data  <- with(test_set,  data_memory(user_index = userId, item_index = movieId, rating = rating))
```

We reate a model object or you can call it a Reference Class object in R by calling the function Reco() from the recosystem package.

In [70]:

```
r <- recosystem::Reco()
```

Recosystem allows to optimize train parameters calling the \$tune option:

In [71]:

```
opts <- r$tune(train_data, opts = list(dim = c(10, 20, 30),
                                       lrate = c(0.1, 0.2),
                                       costp_l2 = c(0.01, 0.1),
                                       costq_l2 = c(0.01, 0.1),
                                       nthread = 6, niter = 10))
```

Finalyy we can train the algorithm with the train_data we prepared before:

In [72]:

```
r$train(train_data, opts = c(opts$min, nthread = 6, niter = 20))
```

iter	tr_rmse	obj
0	0.9684	1.1944e+07
1	0.8729	9.8876e+06
2	0.8380	9.1670e+06
3	0.8153	8.7386e+06
4	0.8002	8.4605e+06
5	0.7888	8.2675e+06
6	0.7794	8.1199e+06
7	0.7713	7.9989e+06
8	0.7645	7.9034e+06
9	0.7584	7.8195e+06
10	0.7533	7.7557e+06
11	0.7486	7.6977e+06
12	0.7444	7.6465e+06
13	0.7406	7.6035e+06
14	0.7371	7.5634e+06
15	0.7341	7.5306e+06
16	0.7311	7.4978e+06
17	0.7284	7.4705e+06
18	0.7259	7.4436e+06
19	0.7236	7.4204e+06

Results

Now that the algorithm is trained we can compute the predicted value using the test_data:

In [74]:

```
y_hat <- r$predict(test_data, out_memory())
```

And verify the RMSE for comparison:

In [76]:

```
RMSE(test_set$rating, y_hat)
```

0.782329091466384

The **RMSE** is therefore **0.78** wich is far below the target of **0.86**

As espected the algorithm built with the recosystem package perform very well for the goal we set up at the beginning. In fact, we have been successful in reduce the target RMSE by **0.08**.

Conclusion

The easiest try to improve our recommendation system could be enrich the data set used with new and different data for example length of the movie or box offices.

An advance improvement can be obtained combining both content based and collaborative filtering methods:

Rating matrix can be also compressed by a neural network. The so called autoencoder is very similar to the matrix factorization. Deep autoencoders, with multiple hidden layers and nonlinearities are more powerful but harder to train. Neural net can be also used to preprocess item attributes so we can combine content based and collaborative approaches.